

Reliable data foundations. Practical **AI systems**.

I help businesses modernise their data platforms, build reliable pipelines, and develop AI systems connected to real business workflows.

- THE PROBLEM

Your AI ambitions depend on your data foundations.

Business-critical data is often spread across disconnected systems, fragile pipelines, and legacy platforms.

Teams spend time reconciling numbers instead of using them. AI prototypes add another layer of complexity when the underlying data is difficult to access or trust.

The challenge is not adding more tools. It is making the system work together.

- THE APPROACH

Connect the foundation to the business outcome.

1

Build the foundation

Bring operational data into a consistent, usable platform.

2

Make it dependable

Establish orchestration, quality checks, clear access, and operational ownership.

3

Put it to work

Enable analytics, business integrations, and carefully evaluated AI applications.

One connected approach — from source systems to the people and applications that use the data.

Architecture thinking. Hands-on engineering.



I'm Faris Hussain, a Munich-based Data & AI Platform Engineer with experience across energy trading, insurance, finance, and data-platform product development.

My work spans enterprise warehouse modelling, cloud ingestion, workflow orchestration, access governance, and multi-agent AI research.

I connect architecture decisions with implementation — and explain the trade-offs to technical and business stakeholders.

Three ways I can help.

Data platform modernisation

Assess existing architecture, design the target platform, and plan migrations around business priorities.

Reliable data engineering

Build and improve ingestion, transformations, orchestration, and governed data delivery.

AI systems & evaluation

Develop domain-specific agents and integrations, then evaluate their behaviour against clear baselines.

Each engagement starts with a defined problem and agreed acceptance criteria.

Modernise without losing sight of what already works.

A migration should solve a business problem — not simply replace one technology with another. I can help with:

- Current-state architecture assessments and dependency mapping.
- Warehouse and lakehouse design, including Data Vault 2.0.
- Multi-source integration and legacy-to-cloud migration planning.
- Phased implementation with reconciliation and explicit cutover criteria.

Typical deliverables: architecture diagrams, decision records, a migration roadmap, and scoped implementation.

Build pipelines that teams can depend on.

Getting data into storage is only the beginning. Reliable delivery also needs incremental processing, stable schemas, sensible retries, useful alerts, and a clear recovery path.

My work connects APIs and operational systems with cloud storage, catalogues, analytical engines, and downstream consumers.

Typical deliverables: working pipelines, orchestration, validation checks, infrastructure as code, and operational handover.

Evaluate AI before expanding its responsibility.

The right question is not simply: "Can we build an agent?" It is: "Where does an agent help, and how will we know?"

My approach starts with a bounded workflow, explicit tool access, and measurable evaluation. Where appropriate, I compare AI-driven behaviour with simpler rule-based alternatives.

Typical deliverables: a scoped prototype, system integrations, an evaluation framework, and recommendations grounded in observed behaviour.

• SELECTED DELIVERY

Finance data ready for analytics.

CONTEXT

A business needed finance and CRM data available in its cloud analytics environment to support reporting and planning.

MY CONTRIBUTION

Built and productionised an API-to-cloud ingestion pipeline, implemented incremental processing, and validated required source datasets and query access.

DELIVERED

Scheduled finance ingestion and queryable finance/CRM datasets, with stakeholder confirmation that the required queries were usable.

Completed ingestion scope; downstream planning-platform integration is a separate, ongoing phase.

• SELECTED DELIVERY

Architecture clarity for due diligence.

CONTEXT

Stakeholders needed a dependable explanation of the existing data platform — and a clear distinction between current capabilities and future plans.

MY CONTRIBUTION

Led an assessment reconciling repository documentation with read-only cloud discovery. Mapped the current architecture, identified gaps, and prepared traceable stakeholder-facing slides.

DELIVERED

A reviewed architecture baseline and current-to-target comparison, with assumptions and unknowns explicitly identified.

Clear evidence makes architecture decisions easier to discuss and defend.

• SELECTED RESEARCH

Multi-agent energy trading.

RESEARCH QUESTION

How do LLM-driven agents compare with rule-based alternatives in simulated energy-trading workflows?

MY CONTRIBUTION

Built energy-specific buyer and seller agents, data-to-simulation transformations, negotiation workflows, and experiment analytics.

EVALUATION SCOPE

Compared behaviour across baseline and stress scenarios, examining negotiation outcomes, economic measures, and model telemetry.

What this demonstrates: the ability to connect domain modelling, agent implementation, and systematic evaluation.

Research simulation — not a claim of live autonomous trading or commercial returns.

- BEYOND THE DEMO

Engineering beyond the happy path.

My wider experience includes:

Platform operations

Improving ingestion continuity, Airflow logging, notifications, and failure visibility.

Governed access

Contributing to data classification, managed access, and infrastructure-as-code adoption.

Reusable data products

Contributing to metadata-driven warehouse tooling and API integration features.

The goal is not just a working demonstration. It is a system another team can understand, use, and maintain.

- HOW WE'D WORK

A clear process, with visible checkpoints.

1

Discover

Understand the business problem, existing systems, constraints, and stakeholders.

2

Define

Agree the scope, deliverables, architecture, and acceptance criteria.

3

Build

Implement in reviewable increments, making dependencies and trade-offs visible.

4

Validate & hand over

Check the agreed behaviour, document limitations, and establish operational ownership.

Progress is measured against accepted deliverables — not just activity.

• ENGAGEMENT OPTIONS

Start with the engagement your project needs.

Architecture assessment

For teams that need clarity before committing to a migration or platform investment.

Scoped implementation

For teams with a defined pipeline, integration, modernisation, or AI evaluation need.

Technical advisory

For teams that need support with architecture decisions, implementation reviews, or delivery planning.

Scope, timeline, availability, and fees are agreed after an initial discussion.

Proposed engagement options for this deck.

- NEXT STEP

Let's define the right next step.

- Modernising a data platform?
- Connecting fragmented business systems?
- Evaluating where AI can genuinely help?

Start with a conversation about the problem, the constraints, and what a useful outcome would look like.

[Connect on LinkedIn](#) [Explore public work on GitHub](#)